

Haripriya Rao

San Francisco, CA | (503) 374-6531 | haripriyaraov@gmail.com | LinkedIn | GitHub | Portfolio

Work Authorization: Cap-Exempt H-1B (Eligible for Immediate Transfer)

Summary

Applied AI and Full-Stack Founding Engineer with **4.5+ years** across multi-agent orchestration, evaluation frameworks, and production systems. Scaled to **255+ agentic steps** at **99.98% uptime**; Google Gemini Cookbook contributor; Claude Certified Architect (CCA-F), AWS Certified AI Practitioner.

Technical Skills

Frontier AI & Orchestration: AI Agents, Agentic Workflows, LangGraph, LangChain, Multi-Agent Systems, RAG Pipelines, Tool Use, Prompt Guardrails, AI Evaluation Frameworks

Frontier Models & SDKs: Claude, GPT, Anthropic API, Model Context Protocol (MCP)

Languages & Full-Stack: Python, TypeScript, JavaScript (ES6+), React, Next.js, Node.js, Flask

Data, Cloud & DevOps: PostgreSQL, Supabase, Redis, ClickHouse, AWS, Vercel, Render, Git, CI/CD

Experience

JPMorgan Chase & Co. (via Think41)

Senior AI Consultant

Remote

Jul 2026 – Present

- Accelerating LLM extraction latency by **52%** (40.8s to 19.6s) at **95%** field agreement by deploying GPT-5.5 and creating automated A/B golden-dataset regression comparison tooling.
- Resolving **100%** of thread misclassifications (**8/8**) across 5–6 level nested forwards with zero false positives via targeted prompt overrides and adversarial test suites.

Oneterra Technologies

Software Engineer 2 / Founding Engineer

Los Angeles, CA

Jan 2024 – Apr 2026

- Delivered an end-to-end multi-tenant platform in 4 months with zero scope creep and **99.98% uptime**, partnering with **NVIDIA stakeholders** and **city officials** across 6 roadmap phases.
- Engineered a recursive DFS algorithm to extract **100+** React components from Figma design trees, decreasing manual UI scaffolding time by **90%**.
- Slashed LLM inference costs by **96%** (from 200KB to **8KB/call**) by redesigning the transformation pipeline to scale horizontally to **255+** agentic steps.
- Cut per-turn orchestrator latency by **40%** via a **22-node asynchronous state machine** using *asyncio.gather* parallel scheduling to eliminate serial round-trips.
- Lifted relevant-result rate **20%** and halved malformed responses (**50%**) using grounded citations and safe JSON parsing recovery.
- Reduced MTTR by **40%** and sustained **99.9% uptime** by owning on-call incident response and instrumenting ClickHouse real-time telemetry dashboards.

Provenir

Full Stack Engineering Intern

New Jersey, USA

May 2023 – Dec 2023

- Maintained **98%** CI/CD pass rate and lowered credit risk SaaS regression defects by **40%**.

Deloitte USI

Software Engineer 1

Bengaluru, India

Jun 2021 – Jul 2022

- Consolidated **15+** data sources into a GraphQL API for **4,000** users, dropping latency from **2s to 400ms**.

Projects

Google Gemini Cookbook (Open Source) | PR #1088

Dec 2025

- Integrated error-aware exponential backoff for 429 rate limits, real-time cost-tracking, and production-grade deployment diagnostics to improve reliability.

Education

University of Southern California

Master of Science in Computer Science

Los Angeles, CA

May 2024

Visvesvaraya Technological University

Bachelor of Engineering in Computer Science

Bengaluru, India

Aug 2021